

Said Reklaoui

AI Engineer · LLMs, RAG & agents

I build AI systems that run in production.

Barcelona, Spain | linkedin.com/in/said-reklaoui-chouaibi-14b42b33 | github.com/sreklaoui

AI FOCUS

LLMs & Prompt/Context Engineering

Retrieval-Augmented Generation

AI Agents & MCP

Fine-tuning & Local Inference

Evaluation & Guardrails

AI in the SDLC

PROFILE

AI Engineer focused on applied AI in production. For the last three years I have built LLM-powered pipelines, retrieval-augmented generation, agent orchestration over MCP and fine-tuned open models on local inference, alongside AI woven through the software development lifecycle itself. What makes that work hold up is the fifteen years underneath it: shipping AI is a systems problem — evaluation, guardrails, cost, latency, observability and a team that can maintain it — and that is where architecture and engineering leadership earn their keep.

SELECTED AI WORK

- **Activity Generator Service** — Generates learning activities for an education platform, either authored by an LLM from a topic or source document or extracted and converted from activities printed in a PDF. Queue-based batch processing with progress and cancellation, a versioned prompt catalogue keyed by activity template and cognitive phase, schema-validated output, multimodal PDF chunking, per-batch cost estimation and tracing over every LLM call. *LLMs, Structured Outputs, Prompt Engineering, Evaluation, LangChain, LangGraph, NestJS, PostgreSQL, Redis*
- **CodeGuard — AI Code Review Platform (personal project)** — Code review platform that analyses structure, dependencies, impact and risk using specialised agents in parallel plus a verification pass. Served open models (Qwen3-Coder 30B → Qwen3.6-35B MoE, FP8) with vLLM on RunPod serverless GPUs behind an async job API; migrating from Ollama to vLLM took throughput from 78 to 226 tok/s aggregate at 5 concurrent sessions, ran a 128k context on a 96 GB GPU, and cut cold start by baking FP8 weights (image 164 GB → 40 GB). Golden-corpus bench for recall/precision and per-review cost billing, which later moved inference to a price-routed gateway. *vLLM, RunPod, GPU Inference, FP8, Qwen, AI Agents, Evaluation, Python*
- **RAG FAQ Agent on Vertex AI** — Semantic retrieval over an indexed knowledge base combined with LLM generation, grounding every answer in real data to reduce hallucinations. *Python, Vertex AI, LangChain, Vector Search*
- **RAG LLM CV — Conversational Assistant** — Turns a static CV into an interactive assistant answering natural-language questions over embeddings and vector search. *Python, LangChain, Streamlit, Vertex AI*

EXPERIENCE

Oneclick

Barcelona, Spain | Current

AI Engineer · VP of Engineering

- Lead a 20+ person engineering organisation across backend, frontend, infrastructure and AI.
- Leading the adoption of AI agents across the whole software lifecycle, on an internal model-agnostic agent platform.
- Raised engineering standards with mutation testing and quality gates in CI, and led an observability push that cut production error issues by ~75%.
- Hands-on in event-driven systems on AWS (Kinesis, Lambda, SQS/SNS, ECS) for a multi-country platform, including high-throughput sync recovery.
- Set architecture direction for a major API migration and for scaling an AI tutor in production.
- Built the Activity Generator Service: an LLM-powered activity generation platform with versioned prompt catalogues, schema-validated output, multimodal document processing and per-batch cost control.

Stayforlong

Barcelona, Spain | Oct 2019 - 2025

AI Engineer

- Designed and shipped email classification pipelines using ML and NLP, combining XGBoost with fine-tuned LLMs (Gemma 3B/7B).
- Led a core event orchestrator automating key support flows — cancellations, modifications and hotel confirmations — integrating Dixa and Google Cloud.
- Orchestrated intelligent agents over the Model Context Protocol (MCP) to infer customer intent and trigger automated actions through external APIs.
- Fine-tuned and converted Hugging Face models to optimised formats (GGUF) for local inference with Ollama.
- Built internal annotation tooling so operations teams could label and curate training data.
- Automated support workflows through multi-label classification and LLM-assisted response generation.
Python, LLMs, Gemma, Hugging Face, Ollama, MCP, XGBoost, Pandas, GCP, Docker.

Backend Engineer

- Built and maintained backend services in Go and PHP using hexagonal architecture, DDD, and strong testing and observability practices.
- Designed decoupled, event-driven systems with Pub/Sub orchestrating service-to-service communication.
- Managed infrastructure and deployment workflows with Jenkins, Cloud Run and Cloud SQL.
Go, PHP, Google Cloud (Cloud Run, Pub/Sub, Cloud SQL), Jenkins, MySQL, Docker, OpenTelemetry.

Letgo

Barcelona, Spain | Jan 2018 - Oct 2019

Backend Engineer & Tech Lead, Automotive Vertical

- Tech lead for the automotive vertical, coordinating design and implementation across the car marketplace.
- Built PHP microservices with hexagonal architecture and asynchronous projections for cross-system data consistency.
- Integrated Carfax vehicle-history data and built event-driven workflows on AWS Lambda (Go) with DynamoDB.
PHP, Go, AWS Lambda, DynamoDB, MySQL, RabbitMQ, Docker, Microservices.

Waynabox

Barcelona, Spain | May 2017 - Dec 2017

Senior Software Engineer

- Built a Symfony microservice to manage an internal business process.
PHP, Symfony, MySQL, RabbitMQ, Docker.

Cdmon

Barcelona, Spain | Mar 2016 - Mar 2017

Operations Manager

- Led a cross-functional team of four to create and integrate a new business opportunity, reporting directly to the CEO.
PHP, Symfony, MySQL, Docker, Elasticsearch.

Atrapalo

Barcelona, Spain | Apr 2013 - Mar 2016

Senior Software Engineer

- Led software projects end to end, from design through delivery.
PHP, Symfony, MySQL, Docker, Elasticsearch, DDD, TDD.

SKILLS

AI & ML LLMs, RAG, AI Agents, MCP, Context Engineering, LangChain, LangGraph, Fine-tuning, Embeddings, Vector Search, Evaluation, NLP, XGBoost, Hugging Face, Ollama, Vertex AI

Leadership Engineering Leadership, AI Strategy, Software Architecture, Delivery & Execution, Quality & Guardrails, Mentoring

Engineering Python, Go, PHP, Microservices, Event-Driven Systems, DDD, Hexagonal Architecture, TDD, Observability

Cloud & Data Google Cloud, AWS, Cloud Run, Pub/Sub, Kafka, Docker, Kubernetes, MySQL, Redis, Elasticsearch, OpenTelemetry

EDUCATION

Master's in Artificial Intelligence · UNIR · 2024 - 2025

Software Engineering Specialist · Universitat Oberta de Catalunya · 2013

Computer Science Degree · Universitat de Barcelona · 2007 - 2010

LANGUAGES

Spanish (native) · Catalan (native) · English (fluent)